

Beginner Level Task...

Task-1 Iris Flowers Classification ML Project :

This particular ML project is usually referred to as the "hello World" of Machine Learning. The iris flowers dataset contains numeric attributes, and it is perfect for beginners to learn about supervised ML algorithms, mainly how to load and handle data. Also, since this is a small dataset, it can easily fit in memory without requiring special transformations or scaling capabilities.

Dataset: <http://archive.ics.uci.edu/ml/datasets/iris>

Name:Madhav Pandey

1. Importing some important libraries and Packages

```
In [18]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno
print("Necessary packages included successfully!")

Necessary packages included successfully!
```

2.Importing the dataset

```
In [17]: df = pd.read_csv("Iris.csv")
df
```

```
Out[17]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
...
145	6.7	3.2	5.2	2.3	Iris-virginica
146	6.2	2.8	6.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica
150 rows x 5 columns					

3. Data Exploration

```
In [18]: r,c = df.shape
print("Number of rows = ",r)
print("Number of columns = ",c)

Number of rows = 150
Number of columns = 5

In [19]: df.head()
```

```
Out[19]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
...
145	6.7	3.2	5.2	2.3	Iris-virginica
146	6.2	2.8	6.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica
150 rows x 5 columns					

```
In [20]: # to display stats about data
df.describe()
```

```
Out[20]:
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764220	0.763161
min	4.300000	2.200000	1.000000	0.100000
max	7.900000	4.300000	6.900000	2.500000
25%	5.000000	3.000000	4.350000	1.300000
75%	6.500000	3.300000	5.100000	1.800000
max	7.900000	4.300000	6.900000	2.500000

```
In [21]: df.info()
```

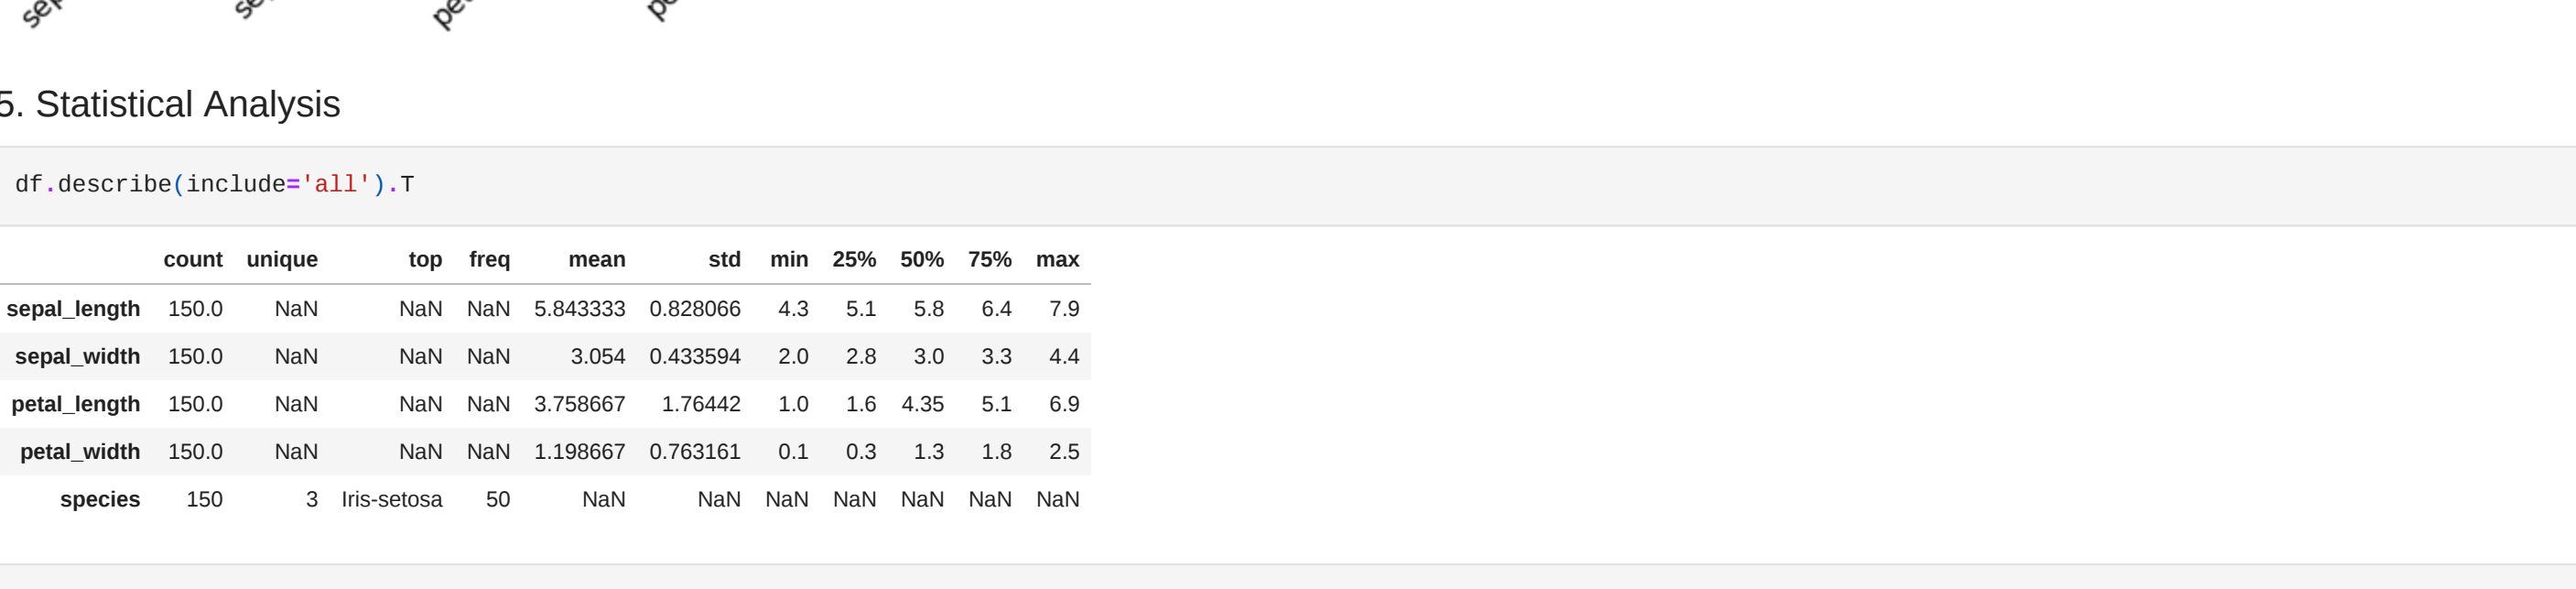
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column             Non-Null Count  Dtype  
---  -
 0   sepal_length       150 non-null    float64
 1   sepal_width        150 non-null    float64
 2   petal_length       150 non-null    float64
 3   petal_width        150 non-null    float64
 4   species            150 non-null    object  
dtypes: float64(4), object(1)
memory usage: 8.0+ KB
```

4. Checking Missing Values

```
In [22]: print("Are there any missing values in the dataset ?", df.isnull().values.any())

Are there any missing values in the dataset ? False
```

```
In [23]: msno.bar(df,figsize=(10,6),color='lightpink')
plt.show()
```



5. Statistical Analysis

```
In [24]: df.describe(include='all').T
```

```
Out[24]:
```


	count	unique	top	freq	mean	std	min	25%	50%	75%	max
sepal_length	150.0	NaN	NaN	NaN	5.843333	0.828066	4.3	5.1	5.8	6.4	7.9
sepal_width	150.0	NaN	NaN	NaN	3.054	0.433594	2.0	2.8	3.0	3.3	4.4
petal_length	150.0	NaN	NaN	NaN	3.758667	1.764220	1.0	1.6	4.35	5.1	6.9
petal_width	150.0	NaN	NaN	NaN	1.198667	0.763161	0.1	0.3	1.3	1.8	2.5
species	150	3	Iris-setosa	50	NaN	NaN	NaN	NaN	NaN	NaN	NaN

```
In [26]: df['species'].unique()
```


```
array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'], dtype=object)
```

6. Parametric Visualization


```
In [27]: g=sns.relplot(x='sepal_length',y='sepal_width',data=df,hue='species',style='species')
g=plt.gcf()
plt.show()
```




```
In [28]: sns.pairplot(df,hue='species')
plt.show()
```




```
In [33]: sns.boxplot(x='species',y='petal_length',data=df)
plt.subplot(2,2,1)
sns.boxplot(x='species',y='petal_width',data=df)
plt.subplot(2,2,2)
sns.boxplot(x='species',y='sepal_length',data=df)
plt.subplot(2,2,3)
sns.boxplot(x='species',y='sepal_width',data=df)
plt.subplot(2,2,4)
sns.boxplot(x='species',y='petal_length',data=df)
plt.show()
```



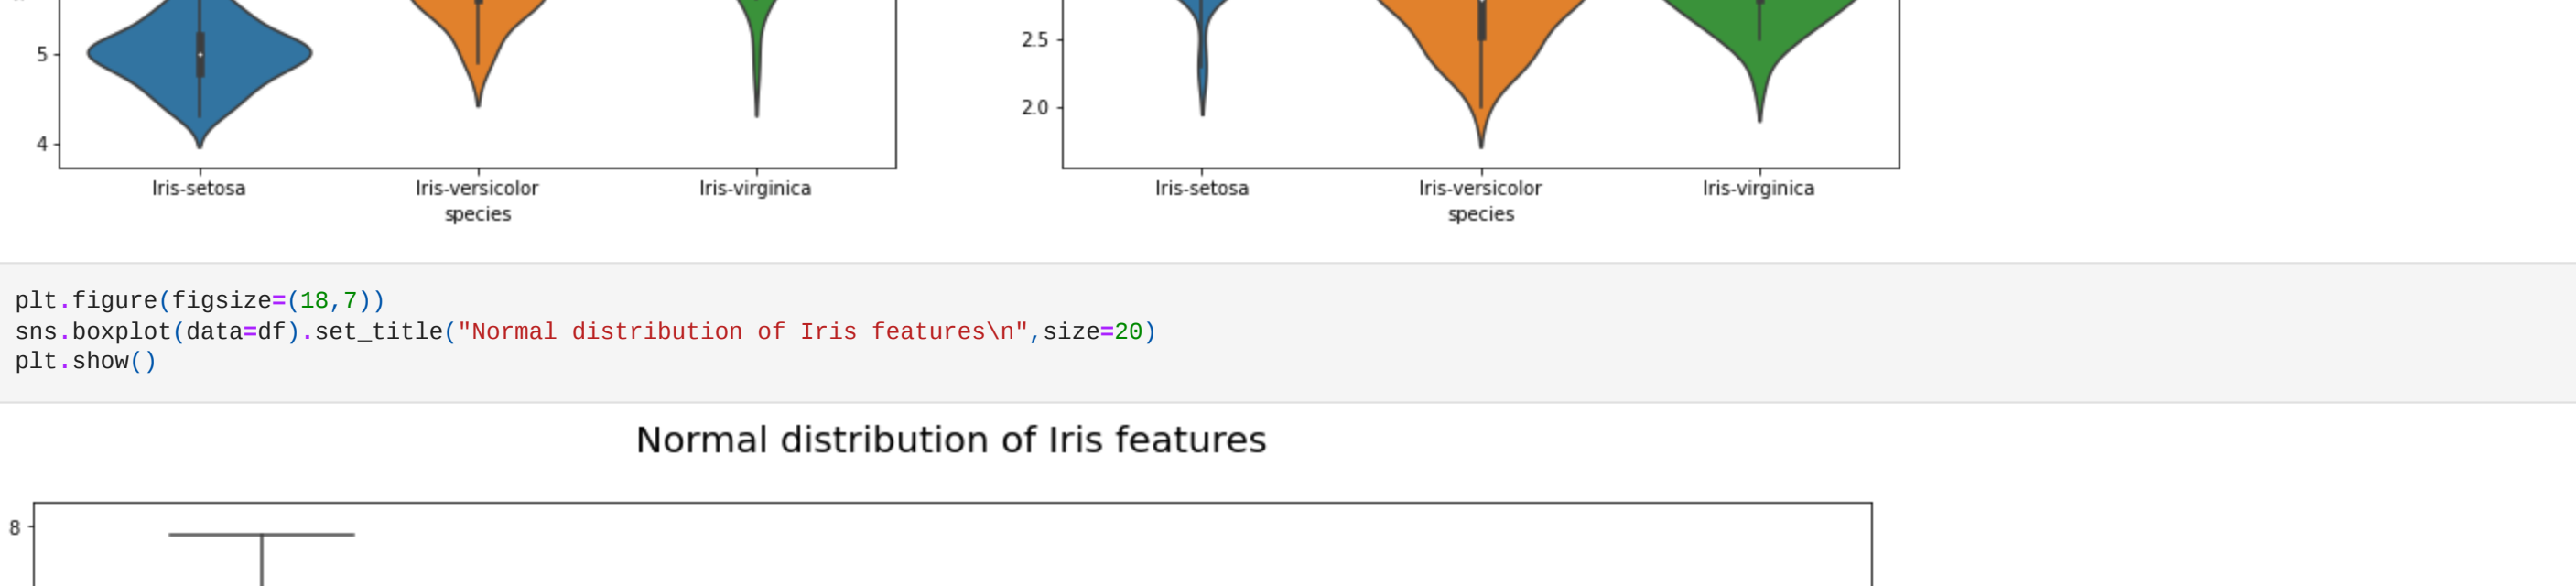
```
In [30]: plt.figure(figsize=(18,10))
plt.subplot(2,2,1)
sns.violinplot(x='species',y='petal_length',data=df)
plt.subplot(2,2,2)
sns.violinplot(x='species',y='petal_width',data=df)
plt.subplot(2,2,3)
sns.violinplot(x='species',y='sepal_length',data=df)
plt.subplot(2,2,4)
sns.violinplot(x='species',y='sepal_width',data=df)
plt.show()
```



```
In [31]: plt.figure(figsize=(18,7))
sns.boxplot(data=df).set_title('Normal distribution of Iris features\n',size=20)
plt.show()
```

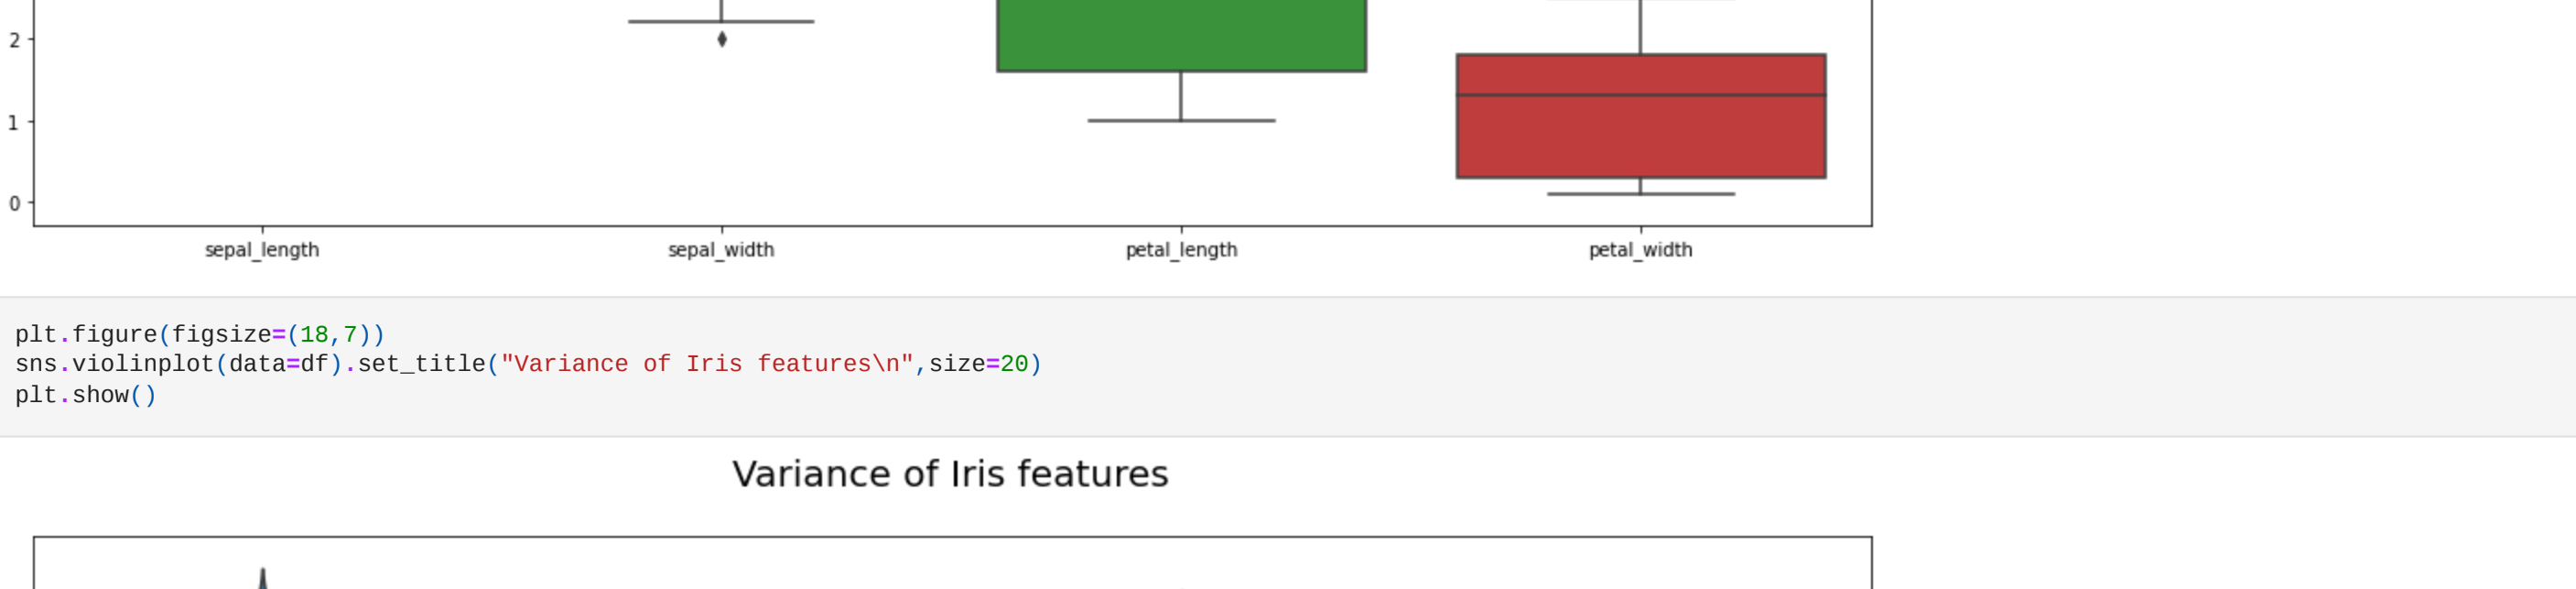


```
In [31]: plt.figure(figsize=(18,7))
sns.violinplot(data=df).set_title('Variance of Iris features\n',size=20)
plt.show()
```



7. Attribute Correlation

```
In [34]: plt.figure(figsize=(16,8))
sns.heatmap(df.corr(),annot=True,fmt='.f',cmap='gist_heat').set_title('Correlation of attributes\n',size=20)
plt.show()
```



```
In [35]: x = df.iloc[:,0:4].values
y = df.iloc[:,4].values
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y = le.fit_transform(y)
```

8. Metric

```
In [36]: from sklearn.metrics import make_scorer,accuracy_score,precision_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score,precision_score,recall_score,f1_score
print("All necessary metrics included!")

All necessary metrics included!
```


9.Model Selection

```
In [37]: from sklearn.model_selection import KFold,train_test_split,cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC, LinearSVC
from sklearn.naive_bayes import GaussianNB
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=0)
print("All Machine Learning packages included!")

All Machine Learning packages included!
```

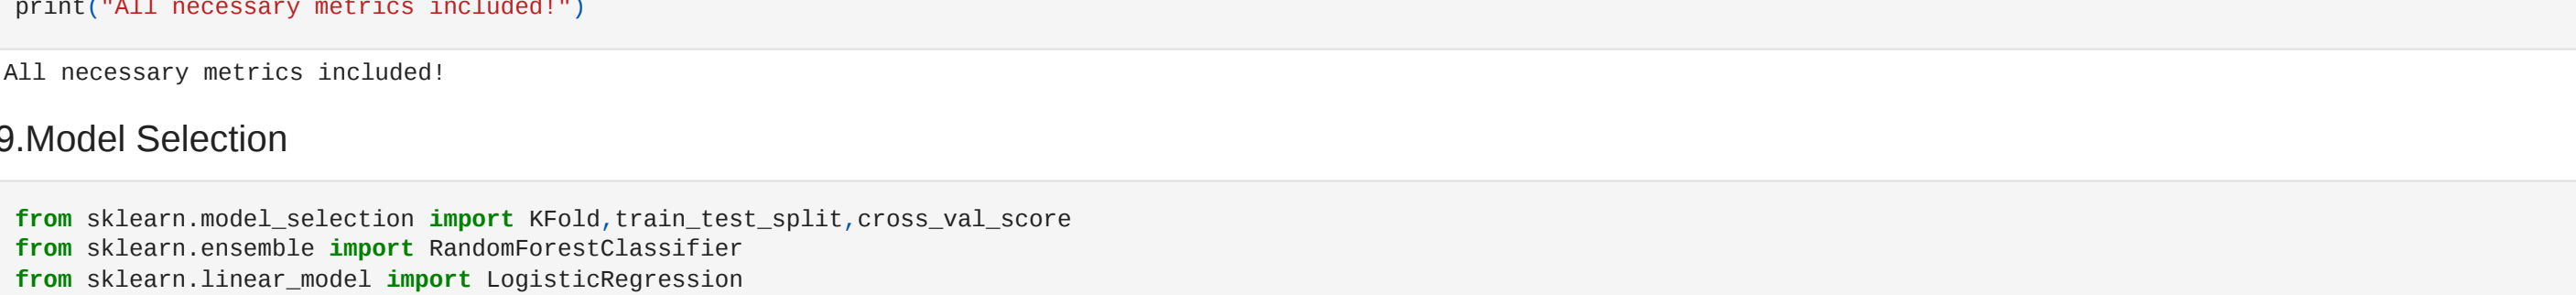
10.Random Forest Rule

```
In [38]: rf = RandomForestClassifier(n_estimators=100)
rf.fit(X_train,y_train)
y_pred = rf.predict(X_test)
rf_acc = round(accuracy_score(y_test,y_pred)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of Random Forest\n",cm)
print("Accuracy of Random Forest = ",acc)
print("Precision of Random Forest = ",prec)
print("Recall of Random Forest = ",recall)
print("F1 score of Random Forest = ",f1)
```




11.Logistic Regression Rule

```
In [39]: lg = LogisticRegression(solver='lbfgs',max_iter=400)
lg.fit(X_train,y_train)
y_pred = lg.predict(X_test)
acc_lg = round(accuracy_score(y_test,y_pred)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of Logistic Regression\n",cm)
print("Accuracy of Logistic Regression = ",acc)
print("Precision of Logistic Regression = ",prec)
print("Recall of Logistic Regression = ",recall)
print("F1 score of Logistic Regression = ",f1)
```



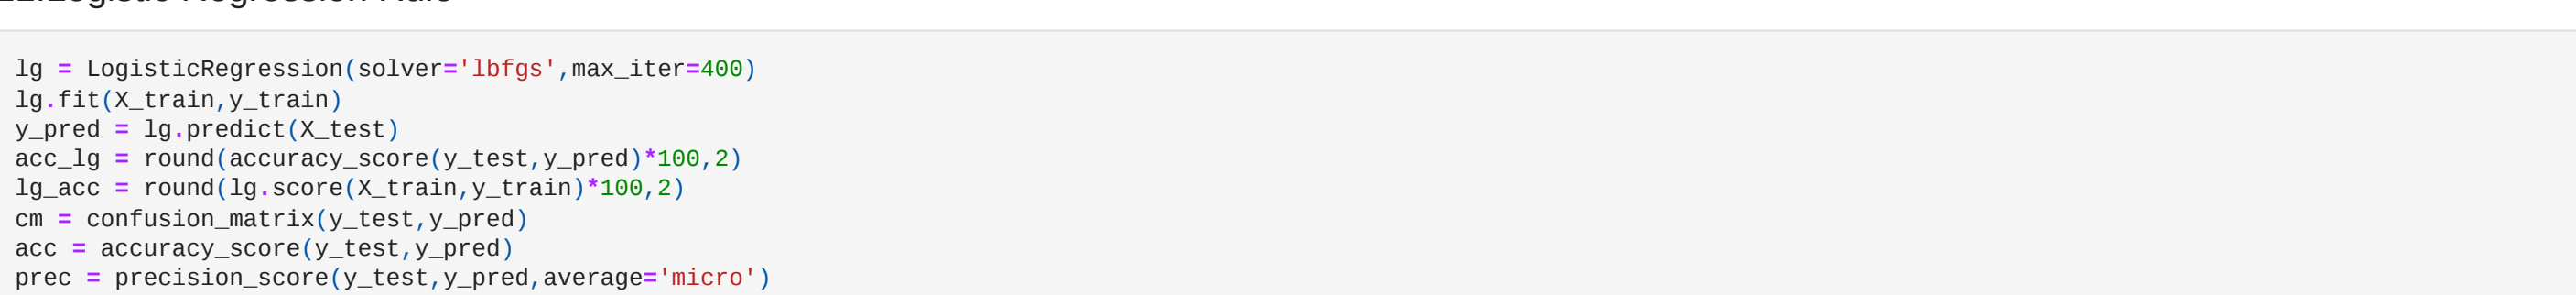
12.K Nearest Neighbours Rule

```
In [40]: knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train,y_train)
y_pred = knn.predict(X_test)
acc_knn = round(accuracy_score(y_test,y_pred)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of K Nearest Neighbour\n",cm)
print("Accuracy of K Nearest Neighbour = ",acc)
print("Precision of K Nearest Neighbour = ",prec)
print("Recall of K Nearest Neighbour = ",recall)
print("F1 score of K Nearest Neighbour = ",f1)
```




13. KNN

```
In [41]: plt.figure(figsize=(20,7))
a_index = list(range(1,50))
a = pd.Series(1,50)
x = range(1,50)
model = KNeighborsClassifier(n_neighbors=1)
model.fit(X_train,y_train)
prediction = model.predict(X_test)
a.append(pd.Series(accuracy_score(y_test,prediction)))
plt.plot(a_index,a,marker='o')
plt.xticks(x)
plt.show()
```



14. Gaussian Naive Bayes rule

```
In [42]: gauss = GaussianNB()
gauss.fit(X_train,y_train)
y_pred = gauss.predict(X_test)
acc_gauss = round(accuracy_score(y_test,y_pred)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of Gaussian Naive Bayes\n",cm)
print("Accuracy of Gaussian Naive Bayes = ",acc)
print("Precision of Gaussian Naive Bayes = ",prec)
print("Recall of Gaussian Naive Bayes = ",recall)
print("F1 score of Gaussian Naive Bayes = ",f1)
```




14. Linear Support Vector Classifier Rule

```
In [43]: lsvc = LinearSVC(max_iter=4000)
lsvc.fit(X_train,y_train)
y_pred = lsvc.predict(X_test)
acc_lsvc = round(accuracy_score(y_test,y_pred)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of Linear Support Vector Classifier\n",cm)
print("Accuracy of Linear Support Vector Classifier = ",acc)
print("Precision of Linear Support Vector Classifier = ",prec)
print("Recall of Linear Support Vector Classifier = ",recall)
print("F1 score of Linear Support Vector Classifier = ",f1)
```



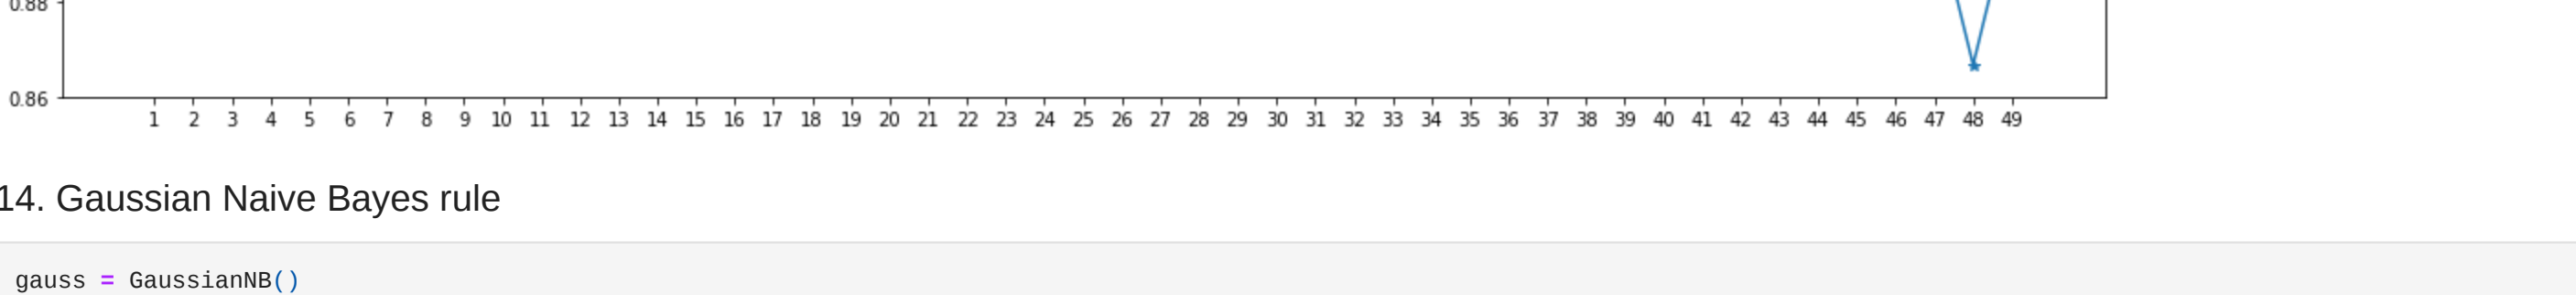
15. Decision Tree Classifier Rule

```
In [44]: dt = DecisionTreeClassifier()
dt.fit(X_train,y_train)
y_pred = dt.predict(X_test)
acc_dt = round(accuracy_score(y_test,y_pred)*100,2)
cm = confusion_matrix(y_test,y_pred)
acc = accuracy_score(y_test,y_pred)
prec = precision_score(y_test,y_pred,average='micro')
recall = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
print("Confusion matrix of Decision Tree\n",cm)
print("Accuracy of Decision Tree = ",acc)
print("Precision of Decision Tree = ",prec)
print("Recall of Decision Tree = ",recall)
print("F1 score of Decision Tree = ",f1)
```



15. Model Scorer Rule

```
In [45]: res = pd.DataFrame({
    'Model':['LR','Logistic Regression','Random Forest','Naive Bayes','Support Vector Regression','Decision Tree'],
    'Score':[acc_knn,acc_lg,acc_rf,acc_gauss,acc_lsvc,acc_dt],
    'Accuracy_score':[knn_acc,lg_acc,rf_acc,gauss_acc,lsvc_acc,dt_acc]
})
res
```



THANK YOU SO MUCH !

```
In [ ]:
```